



## Project 3: Collaboration and Competition

### Table of Contents

1	Author .....	2
2	Acknowledgements.....	2
3	Background and Agent's Learning Environment.....	2
4	Description of the Learning Algorithm and its implementation .....	3
4.1	Introduction .....	3
4.2	DDPG Pseudocode Algorithm Description (Timothy P., et al., 2016) .....	4
4.3	Description of further DDPG improvement.....	5
4.4	Parameter and model .....	5
4.4.1	Agent, Replay Buffer and DDPG.....	5
4.4.2	Model architecture with parameter .....	6
5	Rewards Plot and Results of trained agent.....	8
5.1	Actor NN model and Critic NN model without batch normalisation and with relu activation function.....	8
5.1.1	Reward Plot.....	8
5.1.2	Result of a trained agent.....	8
5.2	Actor NN model and Critic NN model with batch normalisation and with leaky_relu activation function .....	9
5.2.1	Reward Plot.....	9
5.2.2	Result of a trained agent.....	9
5.3	Reward plot and results summary .....	9
6	Ideas for Future Work.....	10
7	References .....	10

# 1 Author

**Name:** [Christian Motschenbacher](#)

**Date:** 03/2019


**Project:** Udacity Deep Reinforcement Learning Nanodegree Program: Project 3 Collaboration and Competition (Tennis)

## 2 Acknowledgements

I would like to thank my wife, who has always encouraged and motivated me at the challenges during my continuous, professional, development courses. Many thanks also to all the teaching staff for the easy and understandable lectures.

## 3 Background and Agent's Learning Environment

The tennis environment has been used in this project. The task of this project was to train two agents, to control rackets to bounce a ball over a net. The following picture show the Tennis environment with two agents.



*Figure 1: Example view of the Tennis Environment*

If an agent hits the ball over the net, it receives a reward of +0.1. If an agent lets a ball hit the ground or hits the ball out of bounds, it receives a reward of -0.01. Thus, the goal of each agent is to keep the ball in play.

The observation space consists of 8 variables corresponding to the position and velocity of the ball and racket. Each agent receives its own, local observation. Two continuous actions are available, corresponding to movement toward or away from the net, and jumping.

The task is episodic, and in order to solve the environment, the agents must get an average score of +0.5 (over 100 consecutive episodes, after taking the maximum over both agents). Specifically,

- After each episode, we add up the rewards that each agent received (without discounting), to get a score for each agent. This yields 2 (potentially different) scores. We then take the maximum of these 2 scores.
- This yields a single score for each episode.

The environment is considered solved, when the average (over 100 episodes) of those scores is at least +0.5.

For the solution of this task it has been used the Unity ML-Agents Reacher Environment from Unity-Technologies, the Python programming language as well the libraries NumPy, PyTorch and others.

The notebook for the training and testing "Collaboration\_and\_Competition\_solution\_training\_testing.ipynb" in this repository contains the chapter "Examine the State and Action Spaces", where you can find an example of the state space and action space as well more detailed information about the environment and project.

The software dependencies installation and configuration of this project are described in the md file in this repository.

## 4 Description of the Learning Algorithm and its implementation

### 4.1 Introduction

The task for this project was to implement a reinforcement algorithm in Python and PyTorch to solve the Tennis environment, which has a continuous observation space and as well a continuous action space. I choose to implement the Deep Deterministic Policy Gradient algorithm (DDPG), which was taught in the lessons. This algorithm is an advancement of the Deep Q Network (DQN) algorithm. This means that the DQN can solve problem with high dimensional observation space as input, whereas those problems can have as output only a discrete or a low-dimensional action space. The model-free, off-policy actor-critic DDPG overcomes this disadvantage and can solve problems with a high dimensional or a continuous observation space as input and a high dimensional or a continuous action space as output. This is possible by using deep function approximators for the observation spaces and as well for the action spaces. One issue of using a non-linear function approximator such as the neural network is that reinforcement learning tends to become unstable or even to diverge. The cause of this instability is for instance the correlation, which is present in the sequence of state, actions, reward and next state. This issue can be improved via the following two features mentioned in the DQN paper and has been implemented in the DDPG.

The first feature is the implementation of a replay memory, which removes some correlations in the observation sequence  $(s_t, a_t, r_t, s_{t+1})$  and act like a smoothing filter. The second feature is to implement a second neural network (target) in addition to the first neural network (local). The target neural network is updated only periodically to reduce the correlation. The implementation of the second neural network made it necessary to parameterize the weights  $\theta_i$  so that the value function changes to  $Q(s, a, \theta_i)$ . The mechanism of replay memory is to store the agent's experiences  $e_t = (s_t, a_t, r_t, s_{t+1})$  at each time-step  $t$  in a data set  $D_t = \{e_1, \dots, e_t\}$ , sample from those uniformly, add those to the observation sequence  $(s_t, a_t, r_t, s_{t+1})$  and use those during learning.

The implemented DDPG algorithm contains all the element mentioned in the paper. Those include replay memory, actor and critic network with batch normalisation, Ornstein-Uhlenbeck noise for exploration. Divergent from the paper, I used as activation function the leaky\_relu instead of relu, which resulted in an additional improvement, which can be seen in the result section of this report.

The following pseudocode Figure 2 show the DDPG, which is used in my project. This DDPG pseudocode is described in the following chapter and as well the additional improvement and research such as replacing the relu activation function with the leaky\_relu, which I have implemented.

---

**Algorithm 1** DDPG algorithm

---

1. Randomly initialize critic network  $Q(s, a|\theta^Q)$  and actor  $\mu(s|\theta^\mu)$  with weights  $\theta^Q$  and  $\theta^\mu$ .
  2. Initialize target network  $Q'$  and  $\mu'$  with weights  $\theta^{Q'} \leftarrow \theta^Q$ ,  $\theta^{\mu'} \leftarrow \theta^\mu$
  3. Initialize replay buffer  $R$
  4. **for** episode = 1, M **do**
  5.   Initialize a random process  $\mathcal{N}$  for action exploration
  6.   Receive initial observation state  $s_1$
  7.   **for** t = 1, T **do**
  8.     Select action  $a_t = \mu(s_t|\theta^\mu) + \mathcal{N}_t$  according to the current policy and exploration noise
  9.     Execute action  $a_t$  and observe reward  $r_t$  and observe new state  $s_{t+1}$
  10.    Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $R$
  11.    Sample a random minibatch of  $N$  transitions  $(s_i, a_i, r_i, s_{i+1})$  from  $R$
  12.    Set  $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}|\theta^{\mu'}))|\theta^{Q'}$
  13.    Update critic by minimizing the loss:  $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$
  14.    Update the actor policy using the sampled policy gradient:
  15.     
$$\nabla_{\theta^\mu} J \approx \frac{1}{N} \sum_i \nabla_a Q(s, a|\theta^Q)|_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s_i}$$
  16.    Update the target networks:
  17.     
$$\theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'}$$
  18.     
$$\theta^{\mu'} \leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}$$
  19.   **end for**
  20. **end for**
- 

Figure 2: Pseudocode of DDPG from DDPG paper

## 4.2 DDPG Pseudocode Algorithm Description ([Timothy P., et al., 2016](#))

The following lines describe the pseudocode of the DDPG in Figure 2 and the implementation of the algorithm.

Line 1: Random initialisation of the weights  $\theta^Q$  for the critic network  $Q$  and weights  $\theta^\mu$  for the actor network  $\mu$ .

Line 2: Initialisation of the target network  $Q'$  and  $\mu'$  by copying the weights from  $Q$  and  $\mu$ .

Line 3: Initialisation of the replay buffer  $R$ .

Line 4 – 20: Episode for loop, which repeats until M has reached.

Line 5: Initialise a random process  $\mathcal{N}$  (noise) for action exploration.

Line 6: Receiving initial observation state  $s_1$ .

Line 7 – 19: Interaction for loop, which repeats until T has reached.

Line 8: Select action according to the current policy and exploration noise.

Line 9: Perform action  $a_t$  and observe reward  $r_t$  as well new state  $s_{t+1}$ .

Line 10: Store the transition  $(s_t, a_t, r_t, s_{t+1})$  into the replay buffer R.

Line 11: Sample a random minibatch of N transitions  $(s_i, a_i, r_i, s_{i+1})$  from replay buffer R.

Line 12: Set target  $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}|\theta^{\mu'}))|_{\theta^{Q'}}$

Line 13: Update the critic by minimizing the loss:  $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$

Line 14: Update the actor policy using the sampled policy gradient:

$$\nabla_{\theta^{\mu}} J \approx \frac{1}{N} \sum_i \nabla_a Q(s, a|\theta^Q)|_{s=s_i, a=\mu(s_i)} \nabla_{\theta^{\mu}} \mu(s|\theta^{\mu})|_{s_i}$$

Line 15: Update the target networks:

$$\begin{aligned} \theta^{Q'} &\leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'} \\ \theta^{\mu'} &\leftarrow \tau \theta^{\mu} + (1 - \tau) \theta^{\mu'} \end{aligned}$$

This pseudocode of the DDPG has been implemented in the three following Python and Pytorch code functions and classes files.

- **ddpg\_agent.py**  
This file contains the class for the agent and the class for the Replay Buffer.
- **model.py**  
This file contains the different neural network models.
- **Collaboration\_and\_Competition\_solution\_training\_testing.ipynb**  
This file contains the code for the training and testing of the agents.

### 4.3 Description of further DDPG improvement.

There are one additional improvements, which has been analysed. This improvement are DDPG model with leaky relu instead of the relu activation function. The models itself, its parameters and as well their performance reward plot are described in in the next chapters.

## 4.4 Parameter and model

### 4.4.1 Agent, Replay Buffer and DDPG

The following parameter are the same for all experiments

#### Environment parameter

- `state_size = 24` # State space of the environment
- `action_size = 2` # Action space of the agent in the environment

### Training parameter

- `n_episodes = 2000`                      # number of training episodes
- `max_t = 1000`                              # maximum number of timesteps per episode

### Agent parameter

- `BUFFER_SIZE = int(1e5)`              # replay buffer size
- `BATCH_SIZE = 256`                      # minibatch size
- `GAMMA = 0.99`                            # discount factor
- `TAU = 1e-3`                                # for soft update of target parameters
- `LR_ACTOR = 1e-4`                        # learning rate of the actor
- `LR_CRITIC = 1e-4`                        # learning rate of the critic
- `WEIGHT_DECAY = 0`                        # L2 weight decay regularization
- `UPDATE_EVERY = 5`                        # when to update the network
- `UPDATE_COUNT = 10`                      # how many times to update the network

## 4.4.2 Model architecture with parameter

The following sections describe the model architecture including their parameter.

### 4.4.2.1 Actor nn model

- Model architecture, parameter

```
class Actor(nn.Module):
    """Actor (Policy) Model."""

    def __init__(self, state_size, action_size, seed, fcl_units=400, fc2_units=300):
        """Initialize parameters and build model.
        Params
        =====
            state_size (int): Dimension of each state
            action_size (int): Dimension of each action
            seed (int): Random seed
            fcl_units (int): Number of nodes in first hidden layer
            fc2_units (int): Number of nodes in second hidden layer
        """
        super(Actor, self).__init__()
        self.seed = torch.manual_seed(seed)
        self.fcl = nn.Linear(state_size, fcl_units)
        self.bnfc1 = nn.BatchNorm1d(num_features=fcl_units)
        self.fc2 = nn.Linear(fcl_units, fc2_units)
        self.bnfc2 = nn.BatchNorm1d(num_features=fc2_units)
        self.fc3 = nn.Linear(fc2_units, action_size)
        self.reset_parameters()

    def reset_parameters(self):
        self.fcl.weight.data.uniform_(*hidden_init(self.fcl))
        self.fc2.weight.data.uniform_(*hidden_init(self.fc2))
        self.fc3.weight.data.uniform_(-3e-3, 3e-3)

    def forward(self, state):
        """Build an actor (policy) network that maps states -> actions."""
        x = F.leaky_relu(self.bnfc1(self.fcl(state)))
        x = F.leaky_relu(self.bnfc2(self.fc2(x)))
        return F.tanh(self.fc3(x))
```

This model has three fully connected layer with the following parameter.

**self.fc1:** 1.fully connected layer with 24 inputs and 400 outputs

**self.bnfc1:** batch normalisation for 1. fully connected layer

**self.fc2:** 2.fully connected layer with 400 inputs and 300 outputs

**self.bnfc2:** batch normalisation for 2. fully connected layer

**self.fc3:** 3.fully connected layer with 300 inputs and 2 outputs

The activation function for fc1 and fc2 is leaky\_relu and it is tanh activation function for fc3.

#### 4.4.2.2 Critic nn model

- Model architecture, parameter

```
class Critic(nn.Module):
    """Critic (Value) Model."""

    def __init__(self, state_size, action_size, seed, fcs1_units=400, fc2_units=300):
        """Initialize parameters and build model.
        Params
        =====
        state_size (int): Dimension of each state
        action_size (int): Dimension of each action
        seed (int): Random seed
        fcs1_units (int): Number of nodes in the first hidden layer
        fc2_units (int): Number of nodes in the second hidden layer
        """
        super(Critic, self).__init__()
        self.seed = torch.manual_seed(seed)
        self.fcs1 = nn.Linear(state_size, fcs1_units)
        self.bnfc1 = nn.BatchNorm1d(num_features=fcs1_units)
        self.fc2 = nn.Linear(fcs1_units+action_size, fc2_units)
        self.bnfc2 = nn.BatchNorm1d(num_features=fc2_units)
        self.fc3 = nn.Linear(fc2_units, 1)
        self.reset_parameters()

    def reset_parameters(self):
        self.fcs1.weight.data.uniform_(*hidden_init(self.fcs1))
        self.fc2.weight.data.uniform_(*hidden_init(self.fc2))
        self.fc3.weight.data.uniform_(-3e-3, 3e-3)

    def forward(self, state, action):
        """Build a critic (value) network that maps (state, action) pairs -> Q-values."""
        xs = F.leaky_relu(self.bnfc1(self.fcs1(state)))
        x = torch.cat((xs, action), dim=1)
        x = F.leaky_relu(self.bnfc2(self.fc2(x)))
        return self.fc3(x)
```

This model has three fully connected layer with the following parameter.

**self.fcs1:** 1.fully connected layer with 24 inputs and 400 outputs

**self.bnfc1:** batch normalisation for 1. fully connected layer

**self.fc2:** 2.fully connected layer with 400 inputs and 300 outputs

**self.bnfc2:** batch normalisation for 2. fully connected layer

**self.fc3:** 3.fully connected layer with 300 inputs and 1 output


The activation function for fc1 and fc2 is leaky\_relu and it is linear activation function for fc3.

## 5 Rewards Plot and Results of trained agent

The following reward plots show the learning reward results of the agents of different network architectures during training.

### 5.1 Actor NN model and Critic NN model without batch normalisation and with relu activation function

#### 5.1.1 Reward Plot



Environment solved in 1936 episodes! Average Score: 0.51


Maximum average score of 0.96 achieved in episode 2000

The above number from the experiment show that this DDPG exceeded an average score of 0.5 over 100 consecutive episodes after 1936 episodes and the highest achieved average score was 0.96. It is quite noticeable that the interaction reward for this particular tennis environment is quite noisy (between about 0 and about 2.6) in comparison to other training environments. It needs to be analysed if this can be improved with prioritized replay memory.

#### 5.1.2 Result of a trained agent

The following run chart shows the reward results of a trained DDPG agent over five episodes.

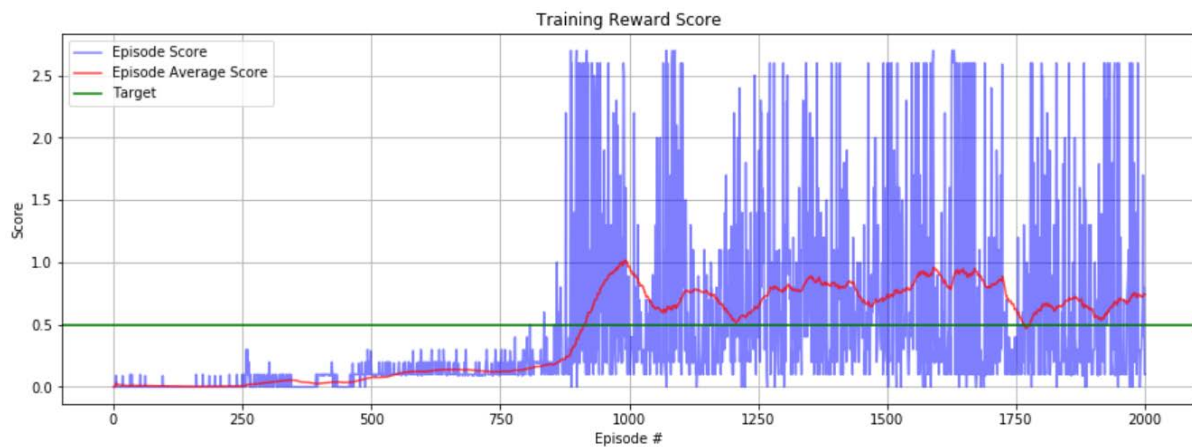
Score of two trained agents: [2.600000038743019, 1.095000016503036, 0.7950000120326877, 2.095000031404197, 0.5450000083073974]





## 5.2 Actor NN model and Critic NN model with batch normalisation and with leaky\_relu activation function

### 5.2.1 Reward Plot



Environment solved in 915 episodes! Average Score: 0.52


Maximum average score of 1.01 achieved in episode 990

The above number from the experiment show that this DDPG exceeded an average score of 0.5 over 100 consecutive episodes after 915 episodes and the highest achieved average score was 1.01. It is quite noticeable that the interaction reward for this particular tennis environment is quite noisy (between about 0 and about 2.6) in comparison to other training environments. It needs to be analysed if this can be improved with prioritized replay memory.

### 5.2.2 Result of a trained agent

The following run chart shows the reward results of a trained DDPG agent over five episodes.

Score of two trained agents: [2.600000038743019, 2.650000039488077, 2.600000038743019, 2.650000039488077, 2.600000038743019]



## 5.3 Reward plot and results summary

The data in the previous section show that the **“actor and critic with batch normalisation and leaky\_relu” perform best** on this task and solve this environment after **915 episodes**. Whereas the **“actor and critic without batch normalisation and relu”**, solves the environment after **1936 episodes**. It is quite noticeable that the interaction reward for this particular tennis environment is quite noisy (between about 0 and about 2.6) in comparison to other training environments. It needs to be analysed if this can be improved with prioritized replay memory.

## 6 Ideas for Future Work

There are several ideas for future work and improvements. Some possible ideas for future work are.

- Ideas and future work to improve the agent's performance
  - Implementing the prioritized experience replay for the DDPG and compare the current results.
  - Implementing and using better RL algorithm such as D4PG
- Some ideas and improvement to speed up the learning and therefore, more time to improve the agents performance are
  - Writing and using a framework, which can be used for an automatized, simplified and faster hyperparameter search.
    - This framework can include for instance and do not be limited to
      - Grid Search
      - Random Search
      - Bayesian Optimisation
    - Run multiple scripts with different parameter in parallel to reduce the hyperparameter search time.
  - Optimise the DDPG algorithm with replay buffer and code by
    - Using vectorised commands where possible instead of for loops
    - Parallelisation and synchronisation of the sampling and learning.
      - One process can sample a few episodes while the other process can perform the learning.

## 7 References

Further resources and references regarding this project and DDPG can be found in the following links.

- Timothy P., et al. "Continuous control with deep reinforcement learning, DDPG (Deep Deterministic Policy Gradients)." <https://arxiv.org/pdf/1509.02971.pdf>
- John Schulman, et al. "PPO (Proximal Policy Optimization Algorithms)." <https://arxiv.org/pdf/1707.06347.pdf>
- Volodymyr Mnih, et al. "Asynchronous Methods for Deep Reinforcement Learning." <https://arxiv.org/pdf/1602.01783.pdf>
- Gabriel Barth-Maroon, et al. "D4PG (Distributed Distributional Deterministic Policy Gradients)." <https://openreview.net/pdf?id=SyZipzbCb>
- Andrej Karpathy. "Andrej Karpathy blog." <http://karpathy.github.io/2016/05/31/rl/>
- Dhruv Parthasarathy. "Write an AI to win at Pong from scratch with Reinforcement Learning." <https://medium.com/@dhrupv/how-to-write-a-neural-network-to-play-pong-from-scratch-956b57d4f6e0>
- Andrej Karpathy, et al. "Evolution Strategies as a Scalable Alternative to Reinforcement Learning." <https://blog.openai.com/evolution-strategies/>
- Shixiang Gu, et al. "Q-Prop: Sample-Efficient Policy Gradient with An Off-Policy Critic." <https://arxiv.org/pdf/1611.02247.pdf>
- John Schulman, et al. "High-Dimensional Continuous Control Using Generalized Advantage Estimation." <https://arxiv.org/pdf/1506.02438.pdf>
- Shixiang Gu, et al. "Continuous Deep Q-Learning with Model-based Acceleration, NAF (normalized advantage functions)." <https://arxiv.org/pdf/1603.00748.pdf>
- Udacity, et al. "Deep Reinforcement Learning Nanodegree." <https://github.com/udacity/deep-reinforcement-learning>